

# Data-Oriented Design



Mike Acton, *Data-Oriented Design and C++*, CppCon 2014

# TLDR

In FP/OOP, you focus on organizing your *code*.

In DOD, you focus on organizing your *data*.

# The “Why” of Data-Oriented

# The Problem

*Picture this: Toward the end of the development cycle, your game crawls, but you don't see any obvious hotspots in the profiler. The culprit? **Random memory access patterns and constant cache misses.** In an attempt to improve performance, you try to parallelize parts of the code, but it takes heroic efforts, and, in the end, you barely get much of a speed-up due to all the synchronization you had to add. To top it off, the code is so complex that fixing bugs creates more problems, and the thought of adding new features is discarded right away. Sound familiar?*

—Noel Llopis, *Data-Oriented Design (Or Why You Might Be Shooting Yourself in The Foot With OOP)*

# The Problem

*That scenario pretty accurately describes almost every game I've been involved with for the last 10 years. The reasons aren't the programming languages we're using, nor the development tools, nor even a lack of discipline. In my experience, **it's object-oriented programming (OOP) and the culture that surrounds it that is in large part to blame** for those problems. OOP could be hindering your project rather than helping it!*

—Noel Llopis, *Data-Oriented Design (Or Why You Might Be Shooting Yourself in The Foot With OOP)*

# The Problem

## Summary

- ▶ Game runs too slowly
- ▶ Why? Too many cache misses
- ▶ Why? Random—rather than sequential and predictable—accesses to memory
- ▶ Why? Common OO usage allocates objects in a tree/graph independent of the memory hierarchy

## The problem with OO culture



# Micro-example

An example of common OO design advice.

*Depend on abstractions, not on concretions*<sup>1</sup>.

Uncle Bob

---

<sup>1</sup>BTW: Concretion is a geological term; it's not the opposite of abstraction.

# Micro-example

```
static private double arrayAvg(double[] it) {  
    double sum = 0.0;  
    for (double x: it) {  
        sum += x;  
    }  
    return sum / it.length;  
}
```

/\* VERSUS \*/

```
static private double iteratorAvg(Iterator<Double> it) {  
    double sum = 0.0;  
    int len = 0;  
    while (it.hasNext()) {  
        sum += it.next();  
        len += 1;  
    }  
    return sum / len;  
}
```

# Micro-example

Method	Time
Array	0.6 sec
Iterator	1.2 sec

- ▶ The “good OO” implementation is twice as slow
- ▶ The array version can be optimized (SIMD, split array + spawn threads)
- ▶ The interface version cannot be optimized.
- ▶ What happens in a system where *everything* follows this advice?

# Micro-example

*Different problems require different solutions.*

*Solving problems you probably don't have creates  
more problems you definitely do have.*

Mike Acton

## The “What” of Data-Oriented

# What is data-oriented design?

*Data-oriented design is an approach to optimising programs by carefully considering the memory layout of data structures*

James McMurray, *An introduction to Data Oriented Design with Rust*

# What is data-oriented design?


*Data-oriented design is an approach to **organizing** programs by carefully considering the memory layout of data structures*

# What is data-oriented design?


*Carefully considering the memory layout of data*




# What is data-oriented design?



# What is data-oriented design?

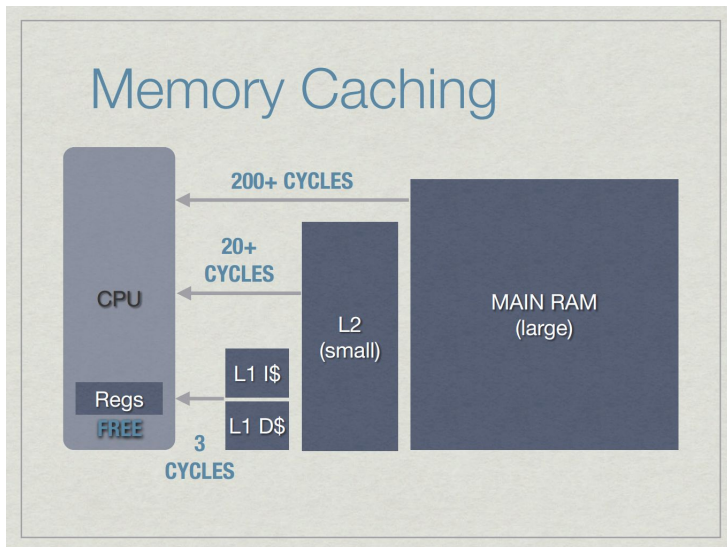


# Computer Memory in Modern Times



Source: <http://gameprogrammingpatterns.com/data-locality.html>


# Computer Memory in Modern Times



Source: Jason Gregory, *Game Engine Architecture*

# Computer Memory in Modern Times

## Demo



## The “How” of Data-Oriented

## Better cache use example

## Cache example

Find the user that logged in most recently



# The input file

```
0:alicea_brittnee:/home/alicea_brittnee:/bin/bash:1576347215
1:carolee_ange:/home/carolee_ange:/bin/bash:1572506918
2:claudie_halsted:/home/claudie_halsted:/bin/bash:1576263841
3:brianne_hizar:/home/brianne_hizar:/bin/bash:1576619633
4:marika_barraza:/home/marika_barraza:/bin/bash:1566035342
5:cristiona_randolf:/home/cristiona_randolf:/bin/bash:1573717947
6:meghan_lief:/home/meghan_lief:/bin/bash:1572370708
7:elly_lemmuela:/home/elly_lemmuela:/bin/bash:1590743916
8:rora_baily:/home/rora_baily:/bin/bash:1567493247
9:sharron_medor:/home/sharron_medor:/bin/bash:1573733695
...
```

# Cache example

Array of Structs (~ Postgresql)

```
type User = struct {  
    id          int64  
    username    string  
    homedir     string  
    shell       string  
    last_login  int64  
}
```

```
latest_timestamp := int64(0)  
latest_username := ""  
for _, user := range users {  
    if user.last_login > latest_timestamp {  
        latest_timestamp = user.last_login  
        latest_username = user.username  
    }  
}
```

# Cache example

## Array of Structs

```
:~) for _ in {1..5}; do go run most_recent_login.go; done
last_login: 1596901628; user: eolanda_pinchas, time: 8.43612ms
last_login: 1596901628; user: eolanda_pinchas, time: 11.060461ms
last_login: 1596901628; user: eolanda_pinchas, time: 11.229049ms
last_login: 1596901628; user: eolanda_pinchas, time: 8.513542ms
last_login: 1596901628; user: eolanda_pinchas, time: 8.479247ms
```

AOS



# Cache example


Hot/Cold Split

```
type UserInfo = struct {  
    id      int64  
    username string  
    homedir  string  
    shell    string  
}  
  
type User = struct {  
    info      *UserInfo  
    last_login int64  
}  
  
latest_timestamp := int64(0)  
latest_username := ""  
for _, user := range users {  
    if user.last_login > latest_timestamp {  
        latest_timestamp = user.last_login  
        latest_username = user.info.username  
    }  
}
```

# Cache example

## Hot/Cold Split

```
:~) for _ in {1..5}; do go run most_recent_login.go; done  
last_login: 1596901628; user: eolanda_pinchas; time: 2.549558ms  
last_login: 1596901628; user: eolanda_pinchas; time: 2.53031ms  
last_login: 1596901628; user: eolanda_pinchas; time: 2.588776ms  
last_login: 1596901628; user: eolanda_pinchas; time: 2.536942ms  
last_login: 1596901628; user: eolanda_pinchas; time: 2.549696ms
```



# Cache example

Struct of Arrays (~ Vertica)


```
type Users = struct {  
    id          []int64  
    username    []string  
    homedir     []string  
    shell       []string  
    last_login  []int64  
}  
  
latest_timestamp := int64(0)  
latest_entity := 0  
for i, last_login := range users.last_login {  
    if last_login > latest_timestamp {  
        latest_timestamp = last_login  
        latest_entity = i  
    }  
}
```

# Cache example

## Struct of Arrays

```
:~ for _ in {1..5}; do go run most_recent_login.go; done  
last_login: 1596901628; user: eolanda_pinchas; time: 1.45387ms  
last_login: 1596901628; user: eolanda_pinchas; time: 1.449812ms  
last_login: 1596901628; user: eolanda_pinchas; time: 1.464322ms  
last_login: 1596901628; user: eolanda_pinchas; time: 1.460788ms  
last_login: 1596901628; user: eolanda_pinchas; time: 1.446045ms
```

SOA



# Cache example

## Takeaways

- ▶ Same complexity ( $O(n)$ ), but in practice constants *do* matter
- ▶ By considering the layout of data, we got a 6x speed-up
- ▶ We made the program 6x faster, not the compiler



An example that has nothing to do with  
cache

# An example that has nothing to do with cache

## The Problem

### segvault-client

Richard found that a lot of time in rtb-gateway was spent calling the function setelement

FUNCTION	CALLS	%
-----	-----	-----
...		
cowboy_handler:execute/2	146	5.01
segvault_protocol:collect/2	129136	7.72
erlang:setelement/3	143367	10.49

# An example that has nothing to do with cache

## The Problem

```
collect([Seg|More], #segvault_response {
    lotame = Lotame,
    adobe = Adobe, ...} = Acc) ->
{SegmentType, SegmentId} = {
    decode_segment_type(Seg bsr ?SEGMENT_BITS),
    Seg band ?SEGMENT_MASK
},
Acc2 = case SegmentType of
    lotame ->
        Acc#segvault_response { lotame = [SegmentId | Lotame] };
    adobe ->
        Acc#segvault_response { adobe = [SegmentId | Adobe] };
```

For those who don't read Erlang:

- ▶ We process one segment at a time
- ▶ We cons the segment to the appropriate list
- ▶ We setelement the new list in the response

# An example that has nothing to do with cache

The Data

```
ADOBE:11111  
ADOBE:11112  
LOTAME:22221  
LOTAME:22222  
LOTAME:22223  
LOTAME:22224  
LIVERAMP3P:33331  
LIVERAMP3P:33332  
LIVERAMP3P:33333  
...
```

Holy fish! The segments are grouped by segment type!

# An example that has nothing to do with cache

## The Solution

- ▶ Accumulate segments in a list (one cons per segment)
- ▶ When the segment type changes, store list in accumulator, empty list (one setelement per *segment type*)

# An example that has nothing to do with cache

The Result



# Extra slides

## Random wisdom from Mike Acton



*The purpose of all programs, and all parts of those programs, is to transform data from one form to another.*

*If you don't understand the data you don't understand the problem.*

*Conversly, understand the problem by understanding the data.*

*Different problems require different solutions.*

*If you have different data, you have a different problem.*

*Where there is one there are many.*

*The more context you have, the better you can make the solution.*

*There is no ideal, abstraction solution to the problem.*

*You can't "future proof"*

*Solve the most common case first, not the most generic.*

# Computer Memory in Modern Times

## 2-minute cache course

- ▶ If a memory address is cached, don't go to memory
- ▶ When fetching from memory, bring an entire cache line (64 bytes) into the cache
- ▶ Best way to take advantage of the cache: keep data contiguous in memory



# Computer Memory in Modern Times

